Distributed Consensus in MongoDB

Spencer T Brody
Senior Software Engineer at MongoDB
@stbrody

Agenda

- Introduction to consensus
- Leader-based replicated state machine
- Elections and data replication in MongoDB
- Improvements coming in MongoDB 3.2

Why use Replication?

- Data redundancy
- High availability

What is Consensus?

- Getting multiple processes/servers to agree on something
- Must handle a wide range of failure modes
  - Disk failure
  - Network partitions
  - Machine freezes
  - Clock skews

Basic consensus

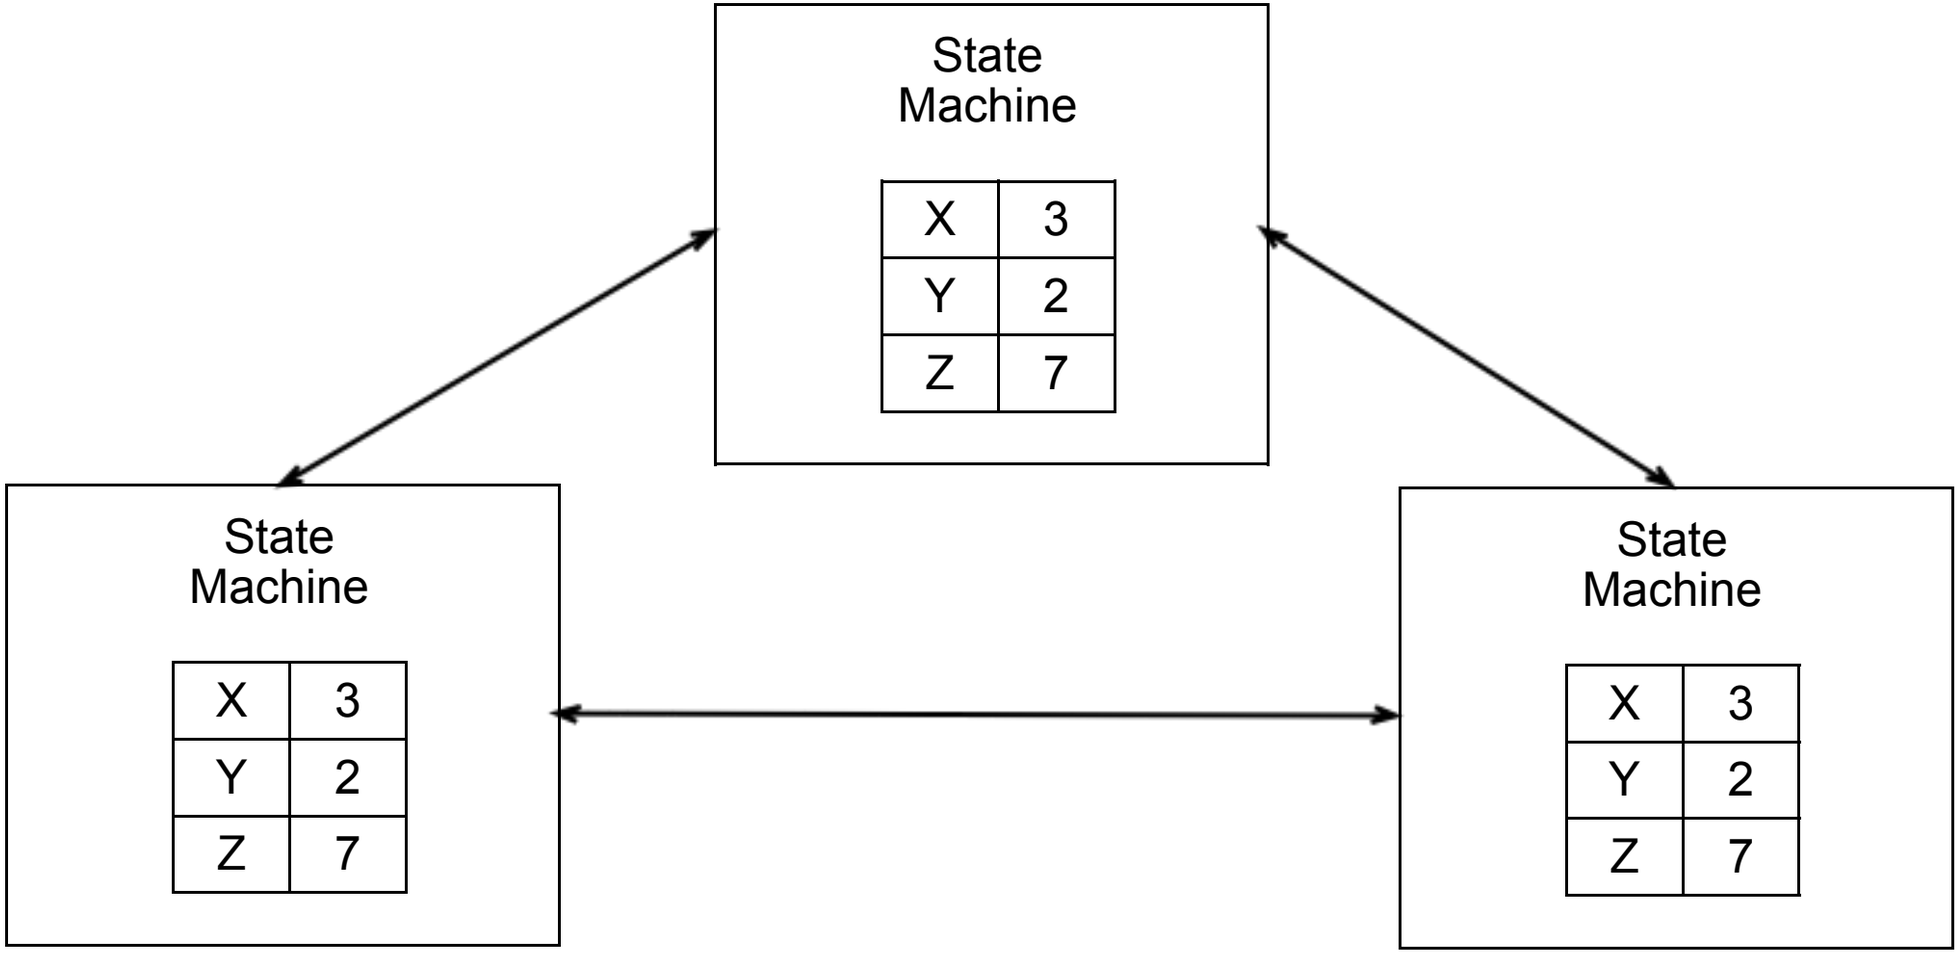

Leader Based Consensus

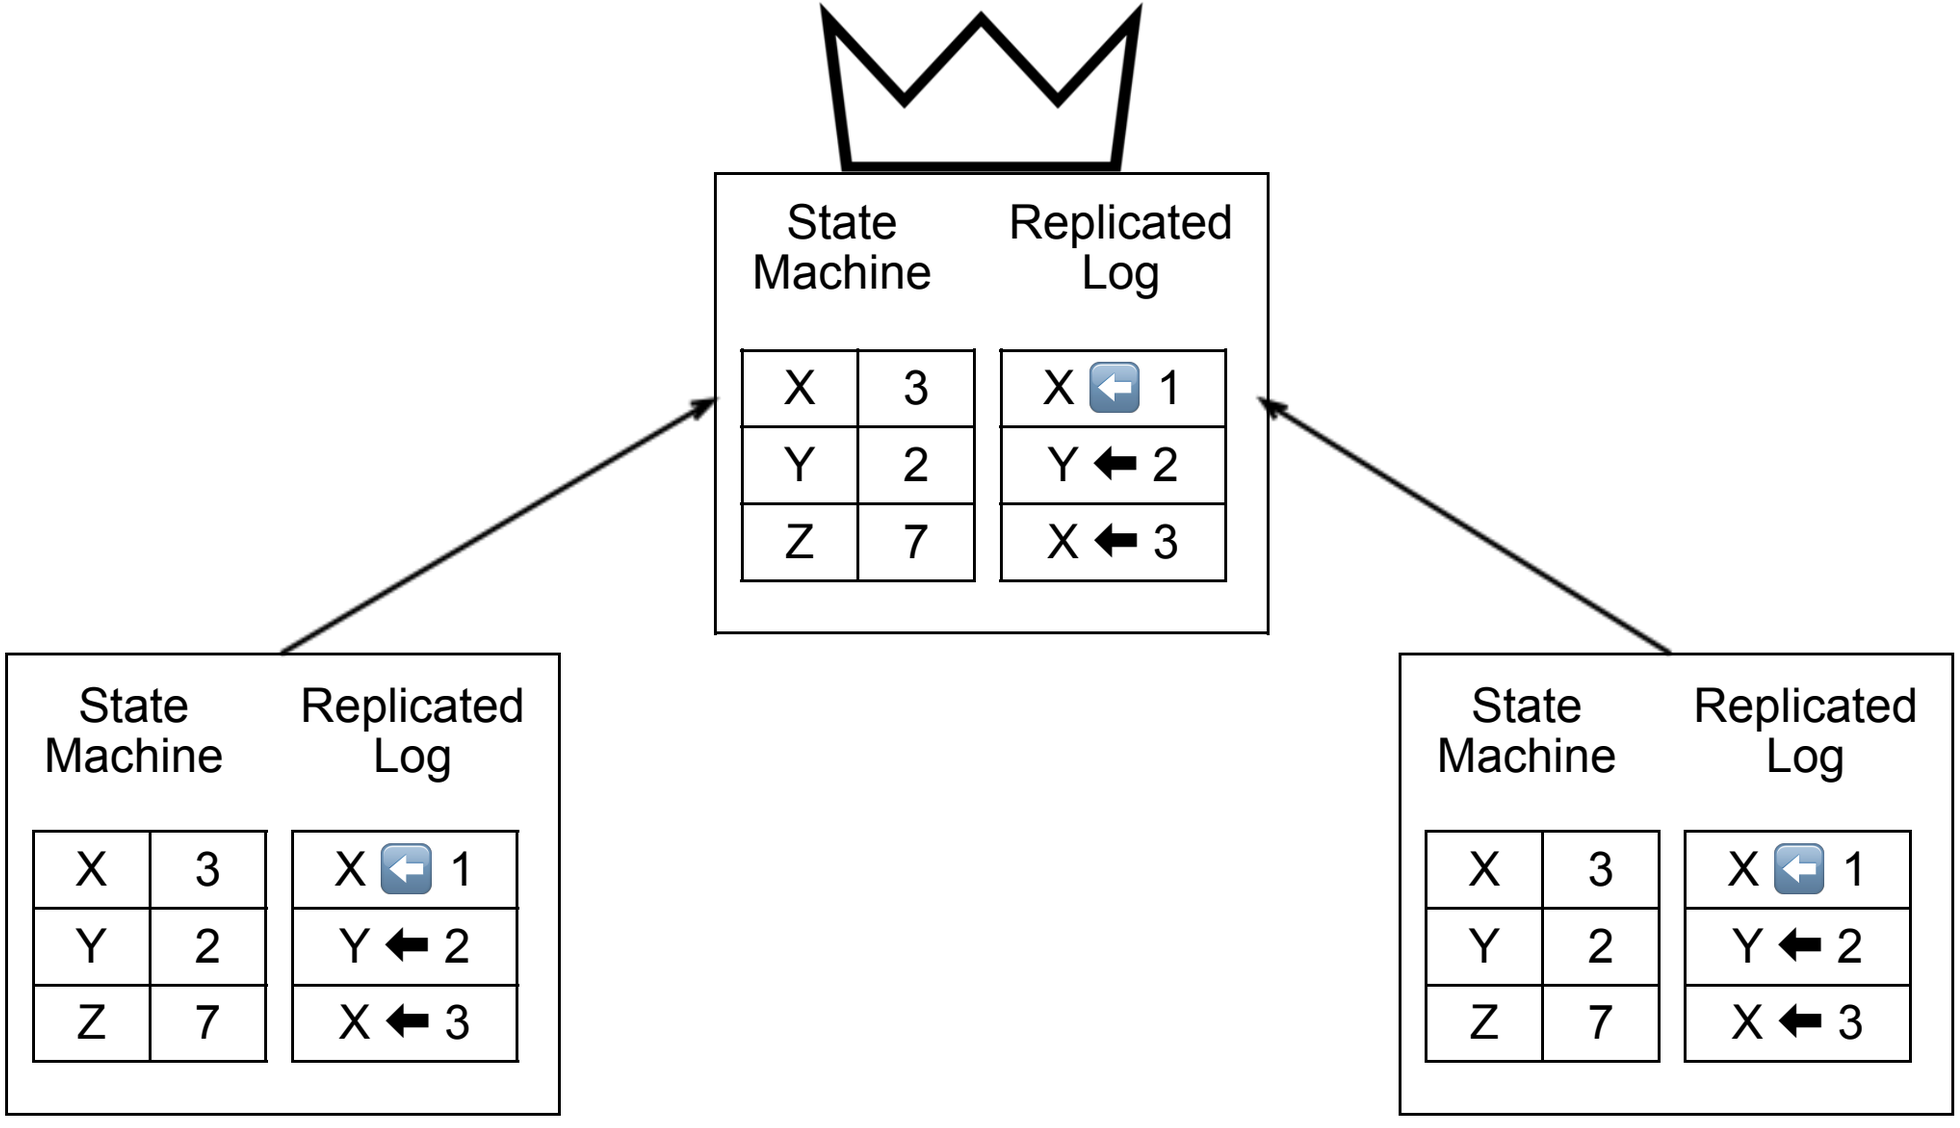

Agenda

- Introduction to consensus
- Leader-based replicated state machine
- Elections and data replication in MongoDB
- Improvements coming in MongoDB 3.2

Elections

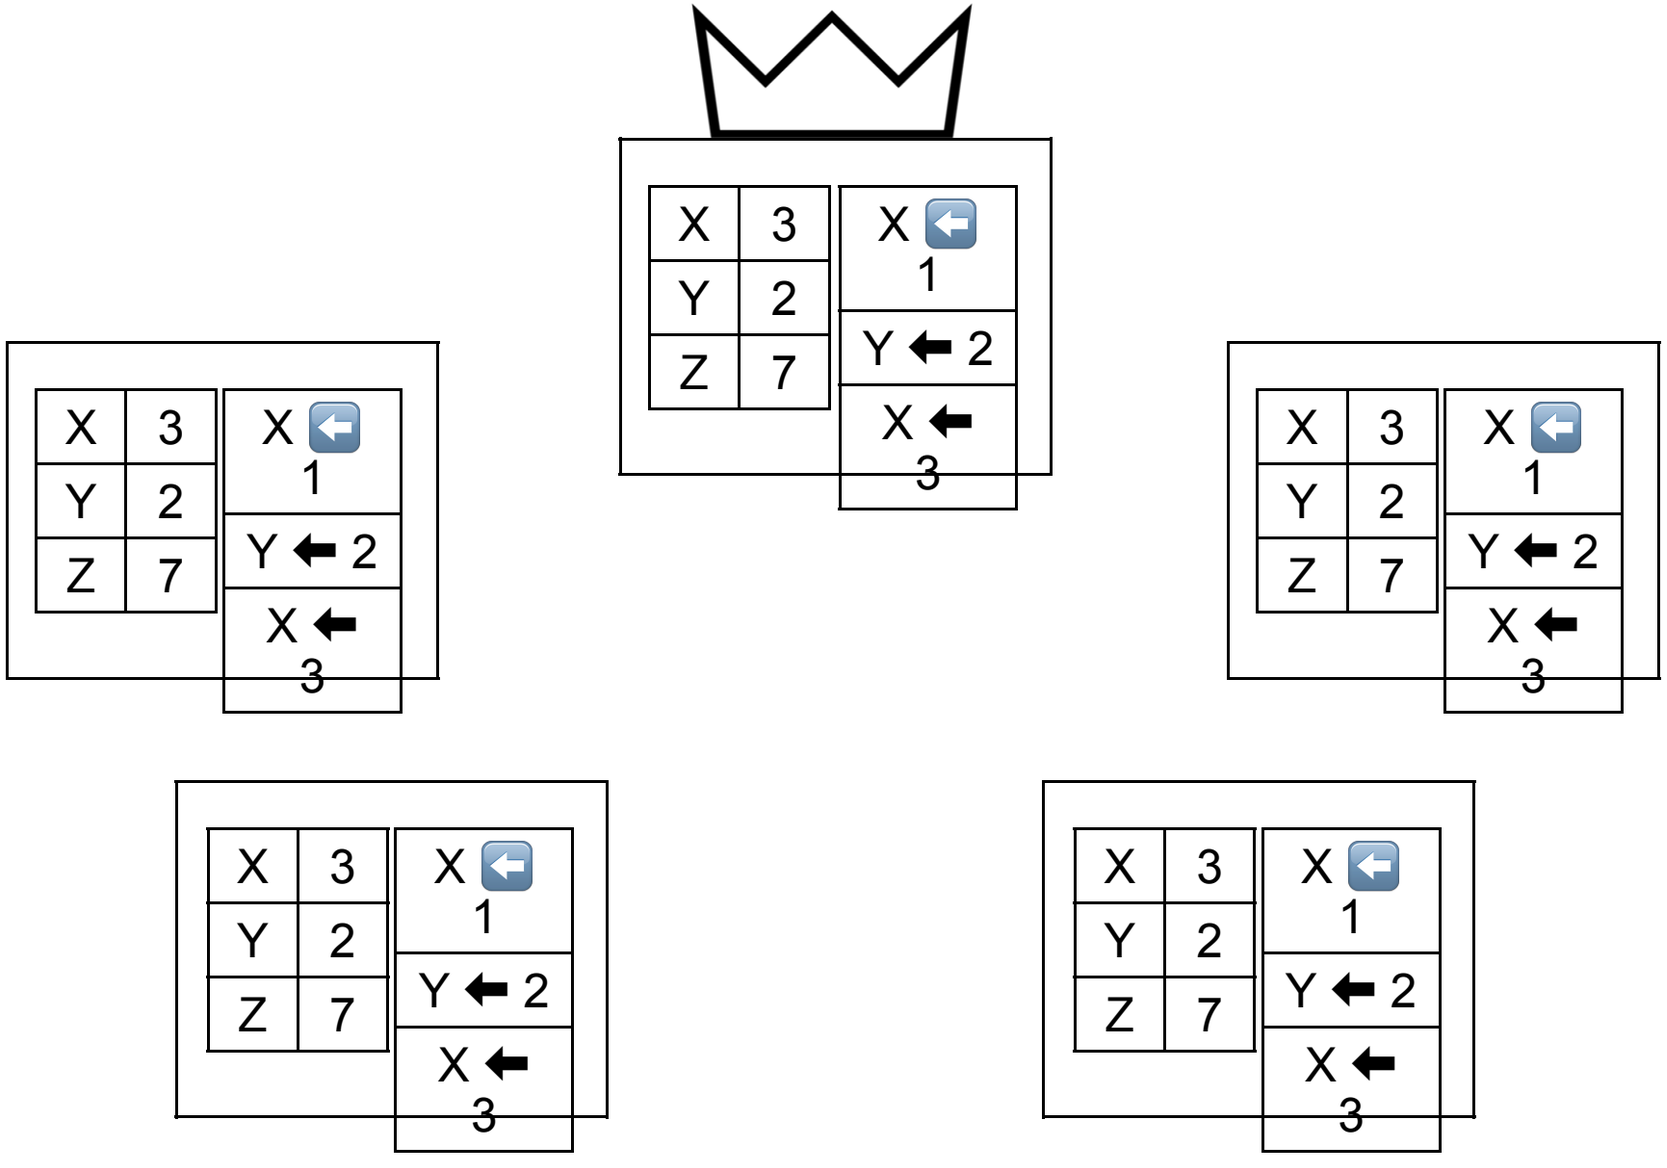

Elections

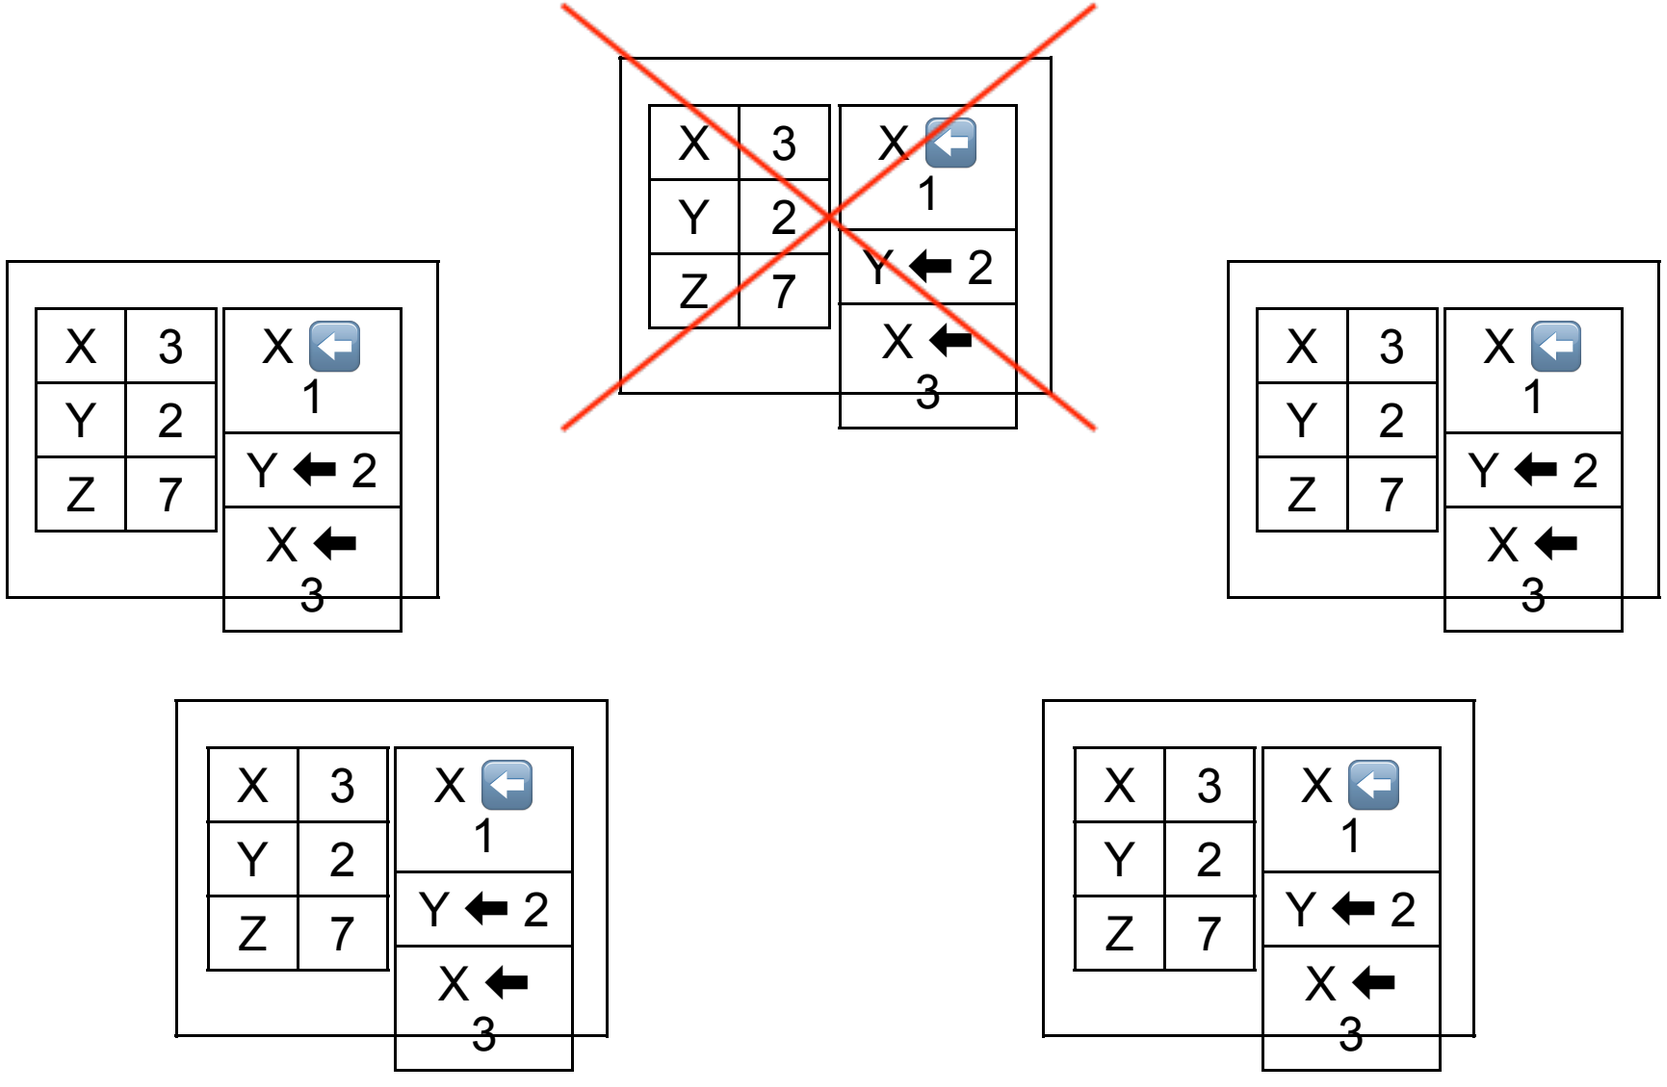

Elections

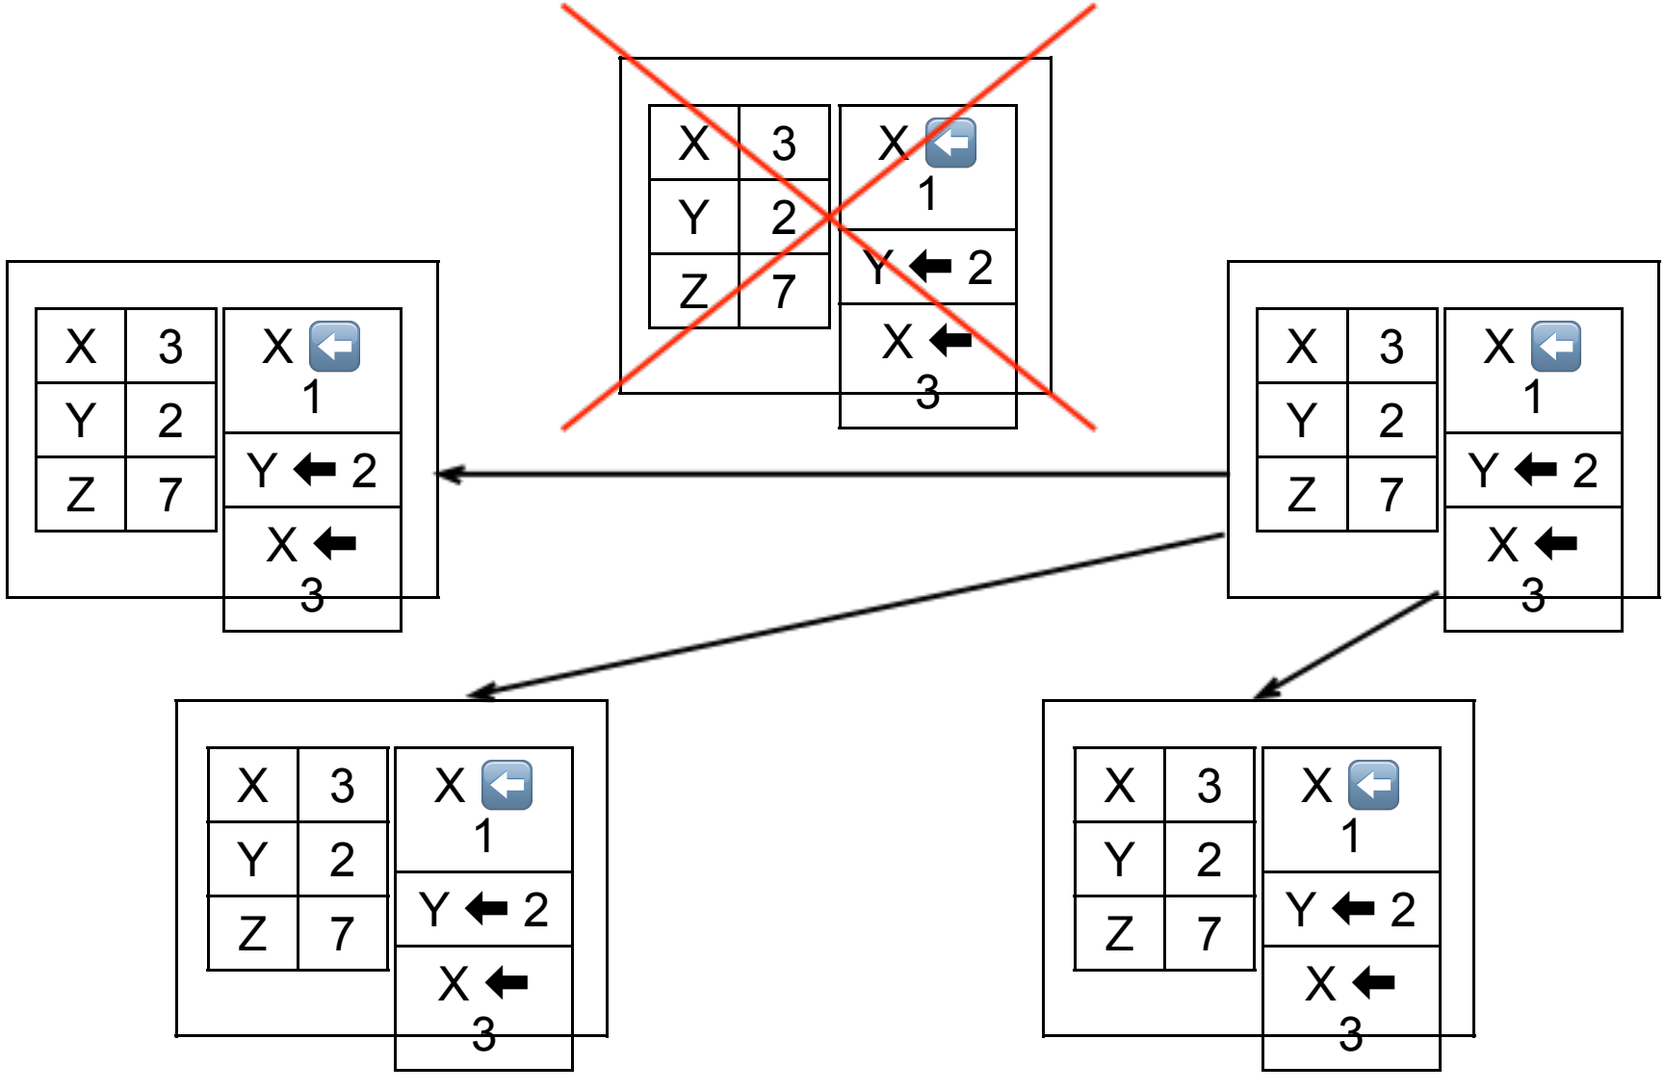

Elections

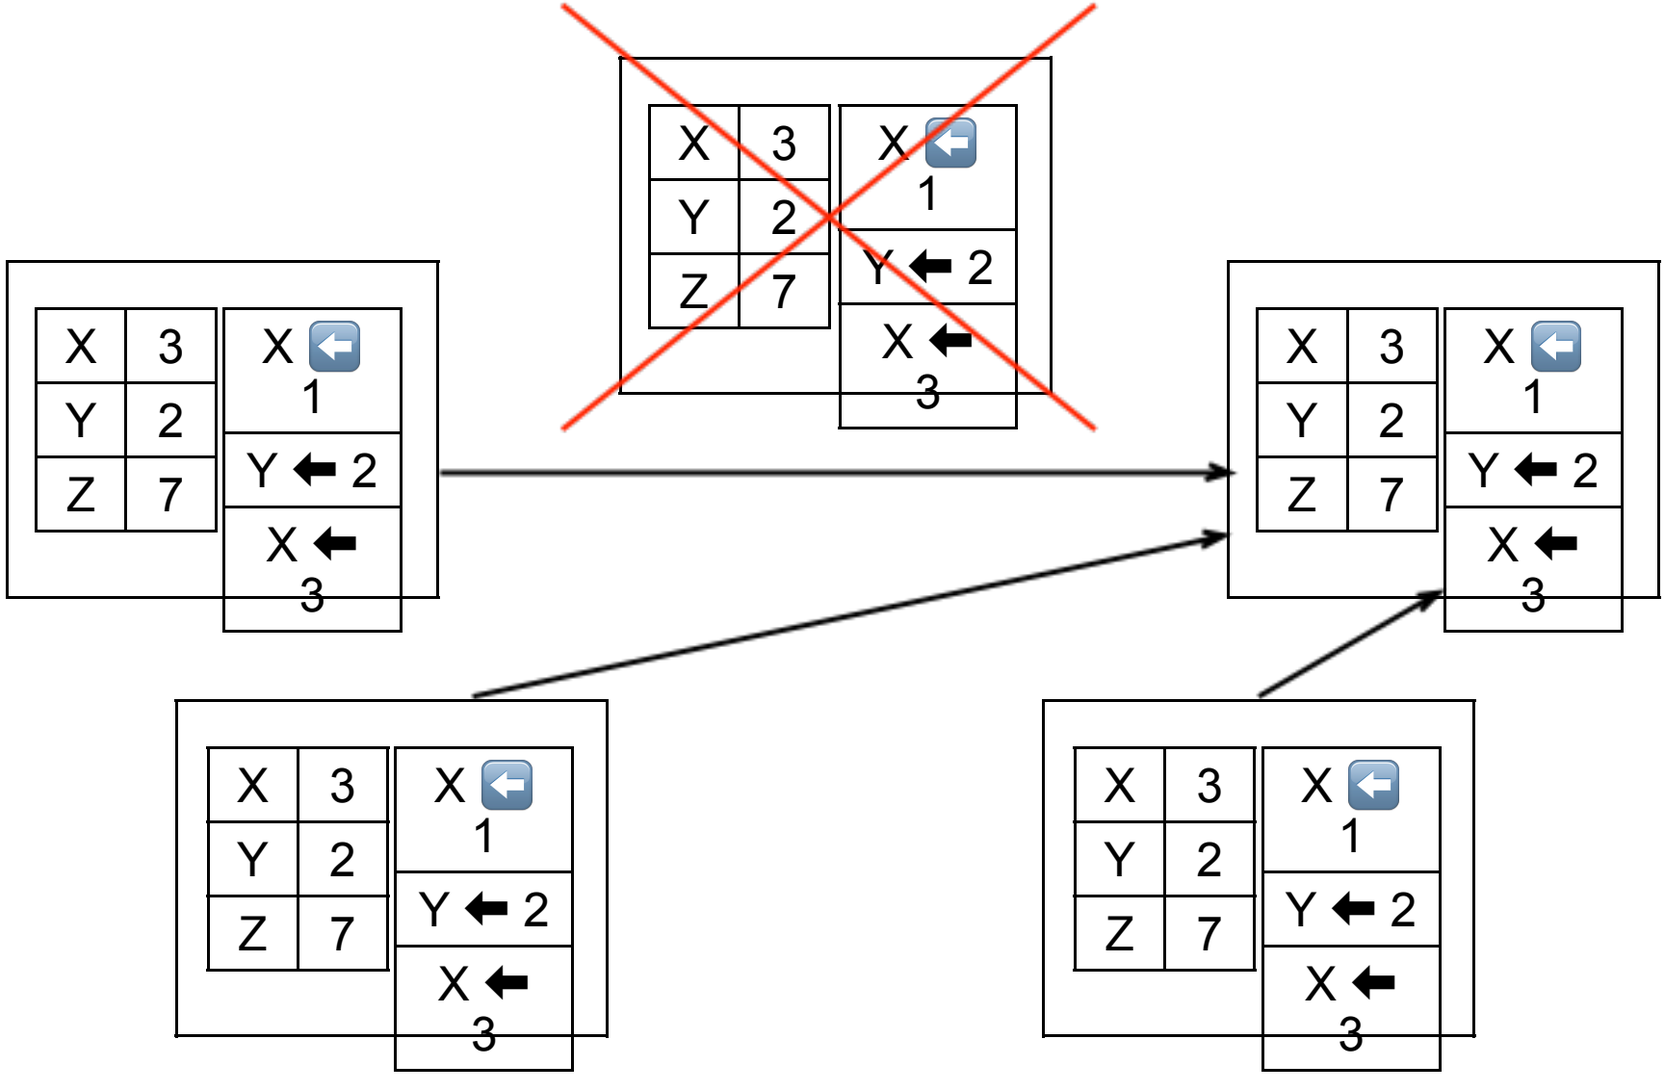

Elections

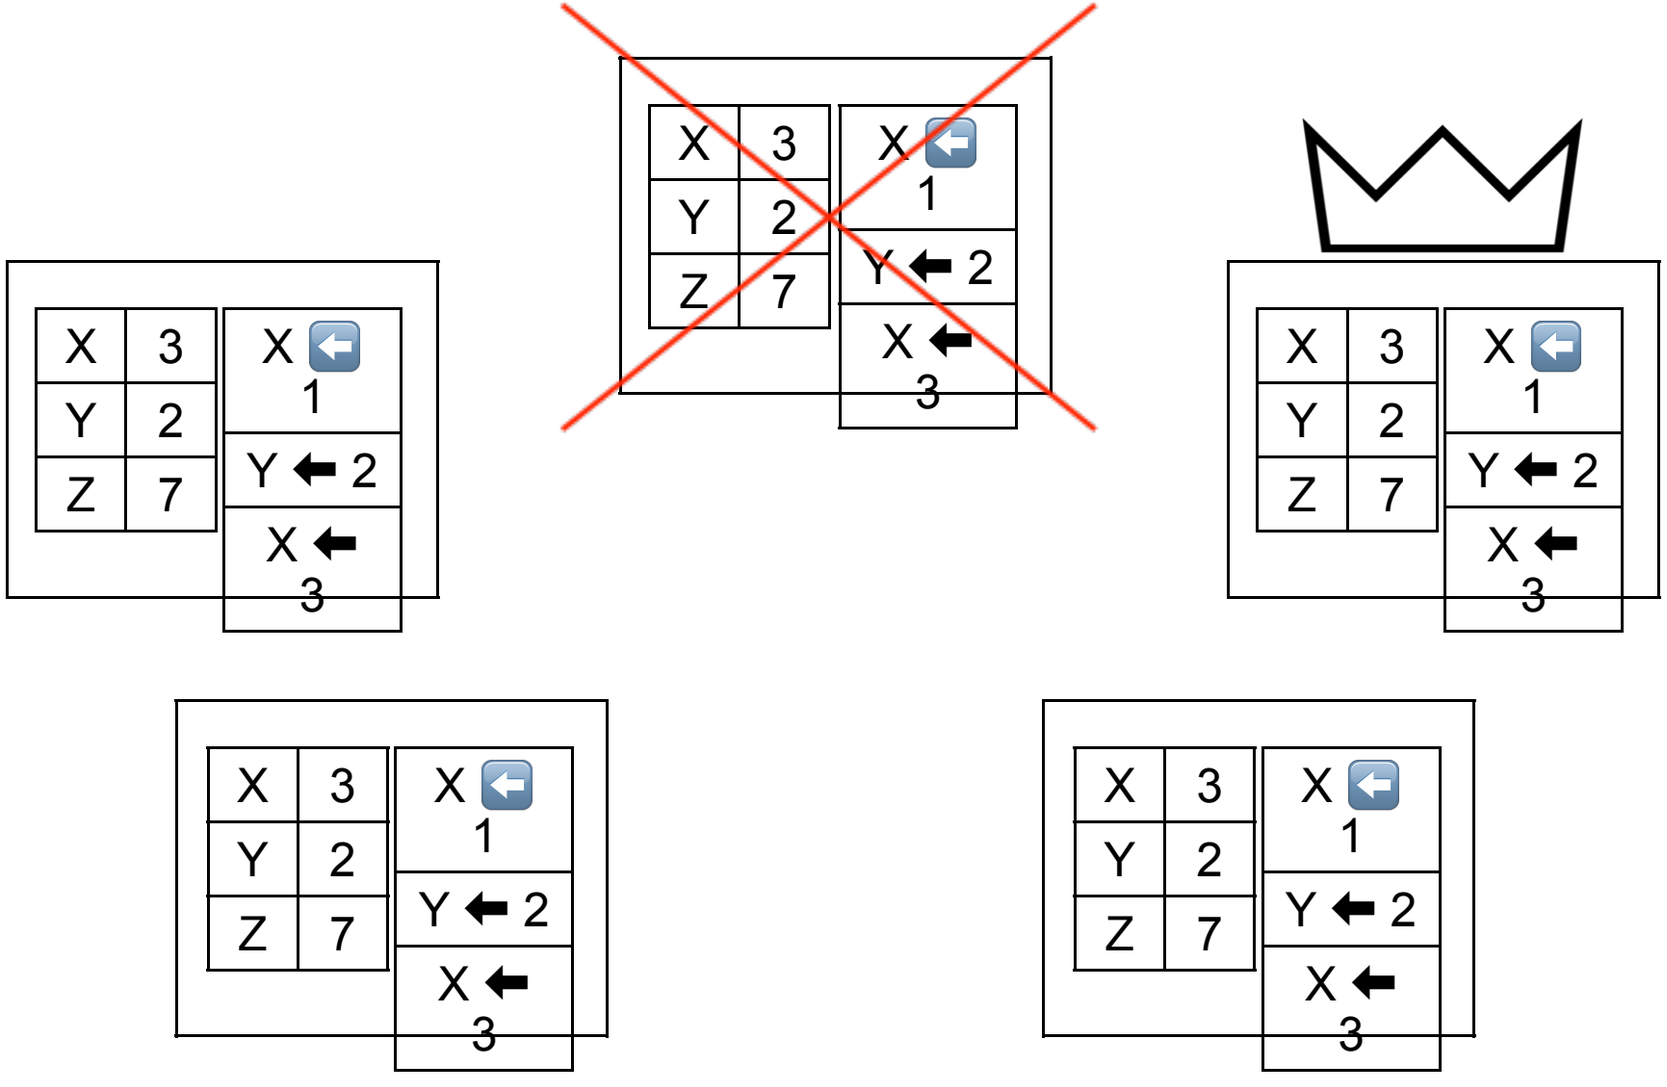

Data Replication

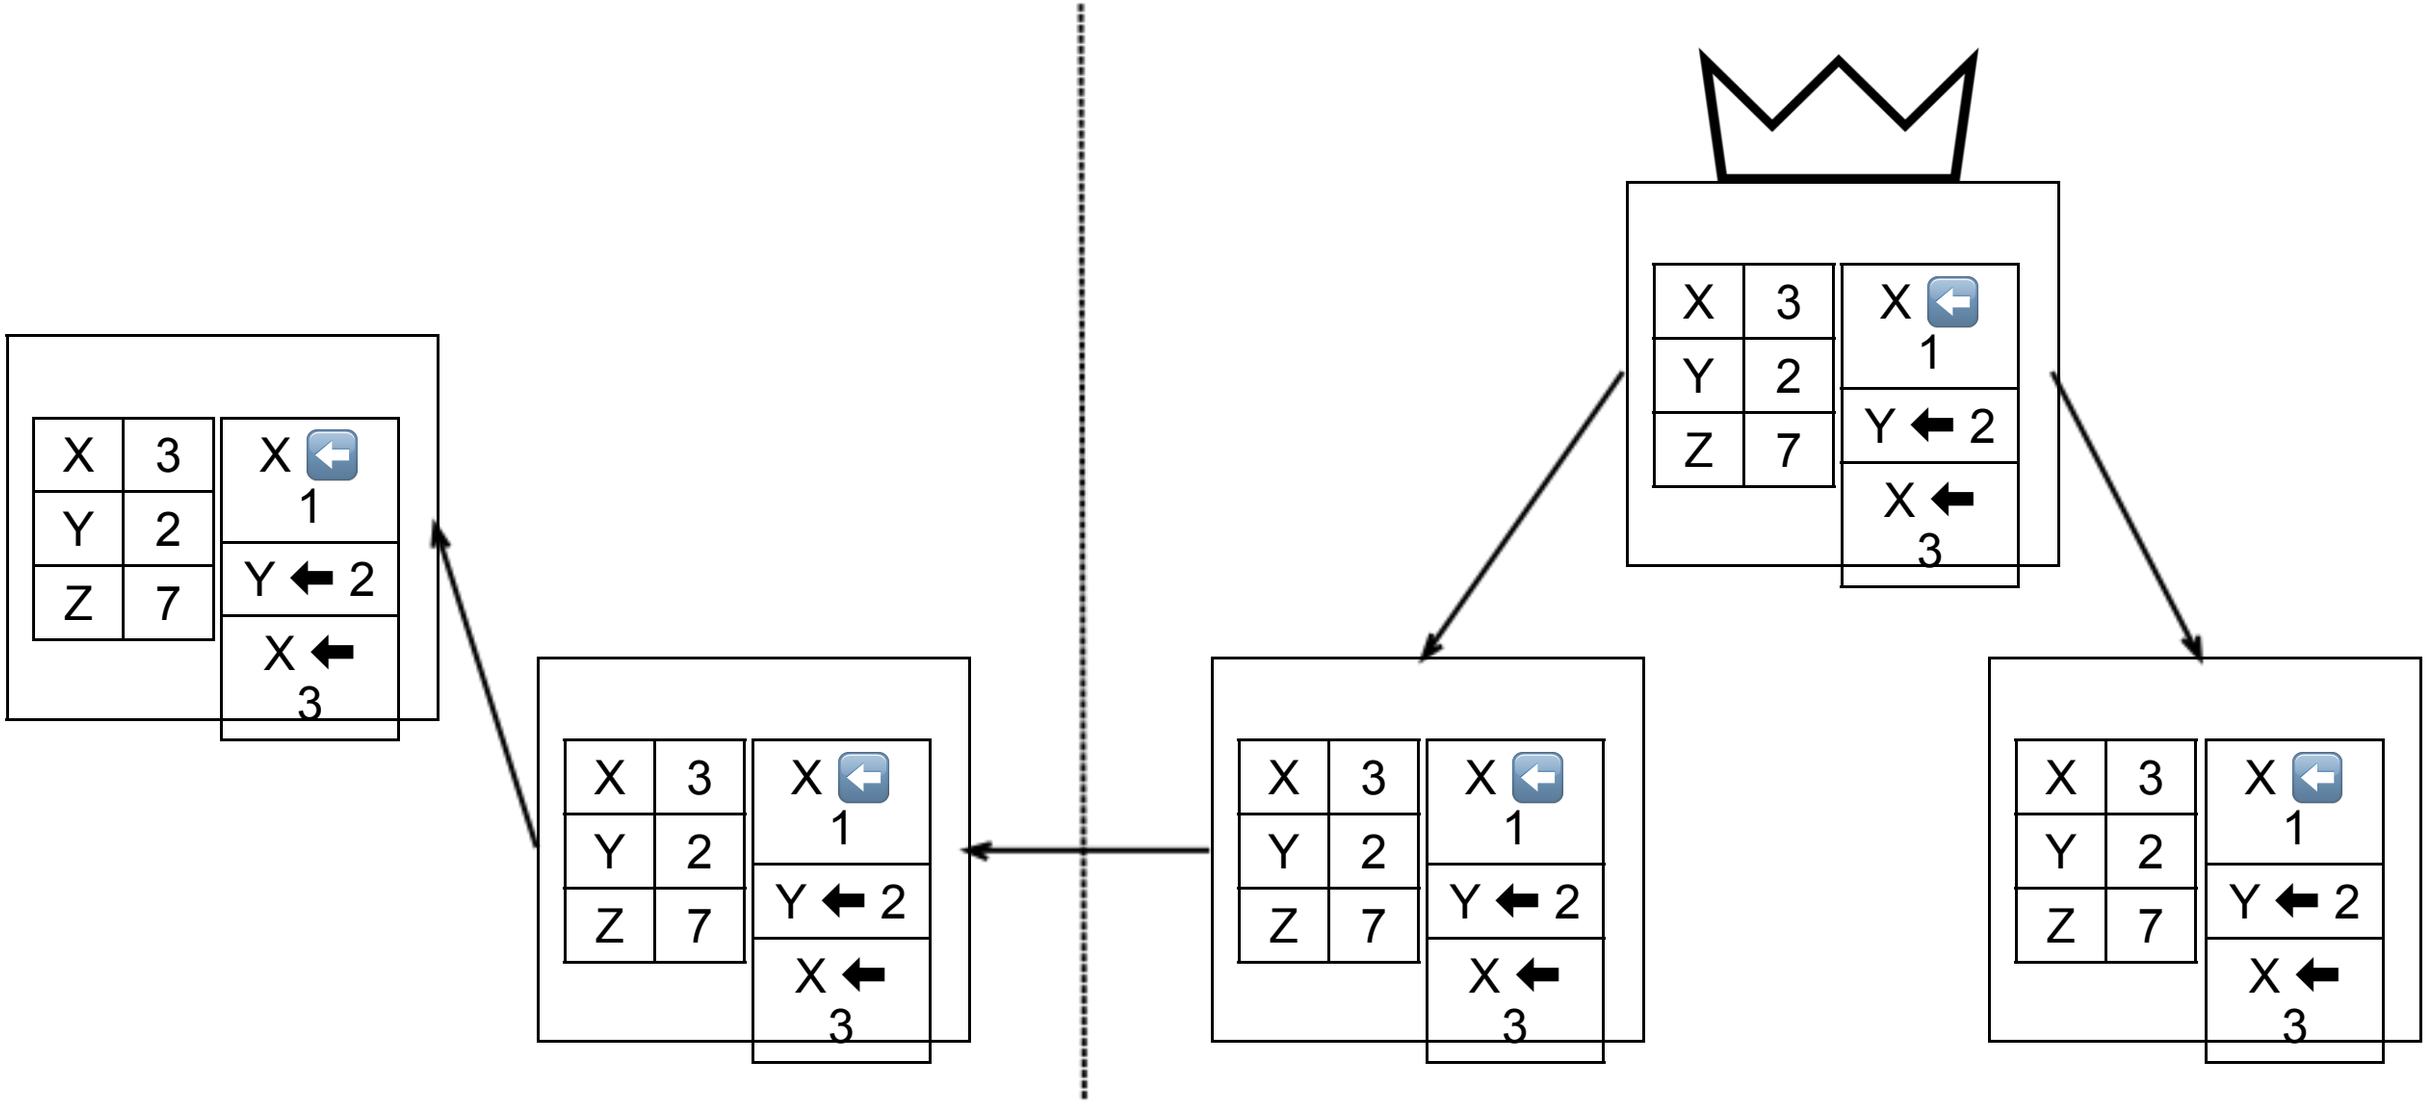

Agenda

- Introduction to consensus
- Leader-based replicated state machine
- Elections and data replication in MongoDB
- Improvements coming in MongoDB 3.2

Agenda

- Introduction to consensus
- Leader-based replicated state machine
- Elections and data replication in MongoDB
- Improvements coming in MongoDB 3.2
  - Goals and inspiration from Raft Consensus Algorithm
  - Preventing double voting
  - Monitoring node status
  - Calling for elections

Goals for MongoDB 3.2

- Decrease failover time
- Speed up detection and resolution of false primary situations

Finding Inspiration in Raft

- “In Search of an Understandable Consensus Algorithm” by Diego Ongaro: https://ramcloud.stanford.edu/raft.pdf
- Designed to address the shortcomings of Paxos
  - Easier to understand
  - Easier to implement in real applications
- Provably correct
- Remarkably similar to what we’re doing already

Raft Concepts

- Term (election) IDs
- Monitoring node status using existing data replication channel
- Asymmetric election timeouts

Preventing Double Voting

- Can’t vote for 2 nodes in the same election
- Pre-3.2: 30 second vote timeout
- Post-3.2: Term IDs
- Term:
  - Monotonically increasing ID
  - Incremented on every election *attempt*
  - Lets voters distinguish elections so they can vote twice quickly in different elections.

Monitoring Node Status

- Pre-3.2: Heartbeats
  - Sent every two seconds from every node to every other node
  - Volume increases quadratically as nodes are added to the replica set
- Post-3.2: Extra metadata sent via existing data replication channel
  - Utilizes chained replication
  - Faster heartbeats = faster elections and detection of false primaries

Data Replication

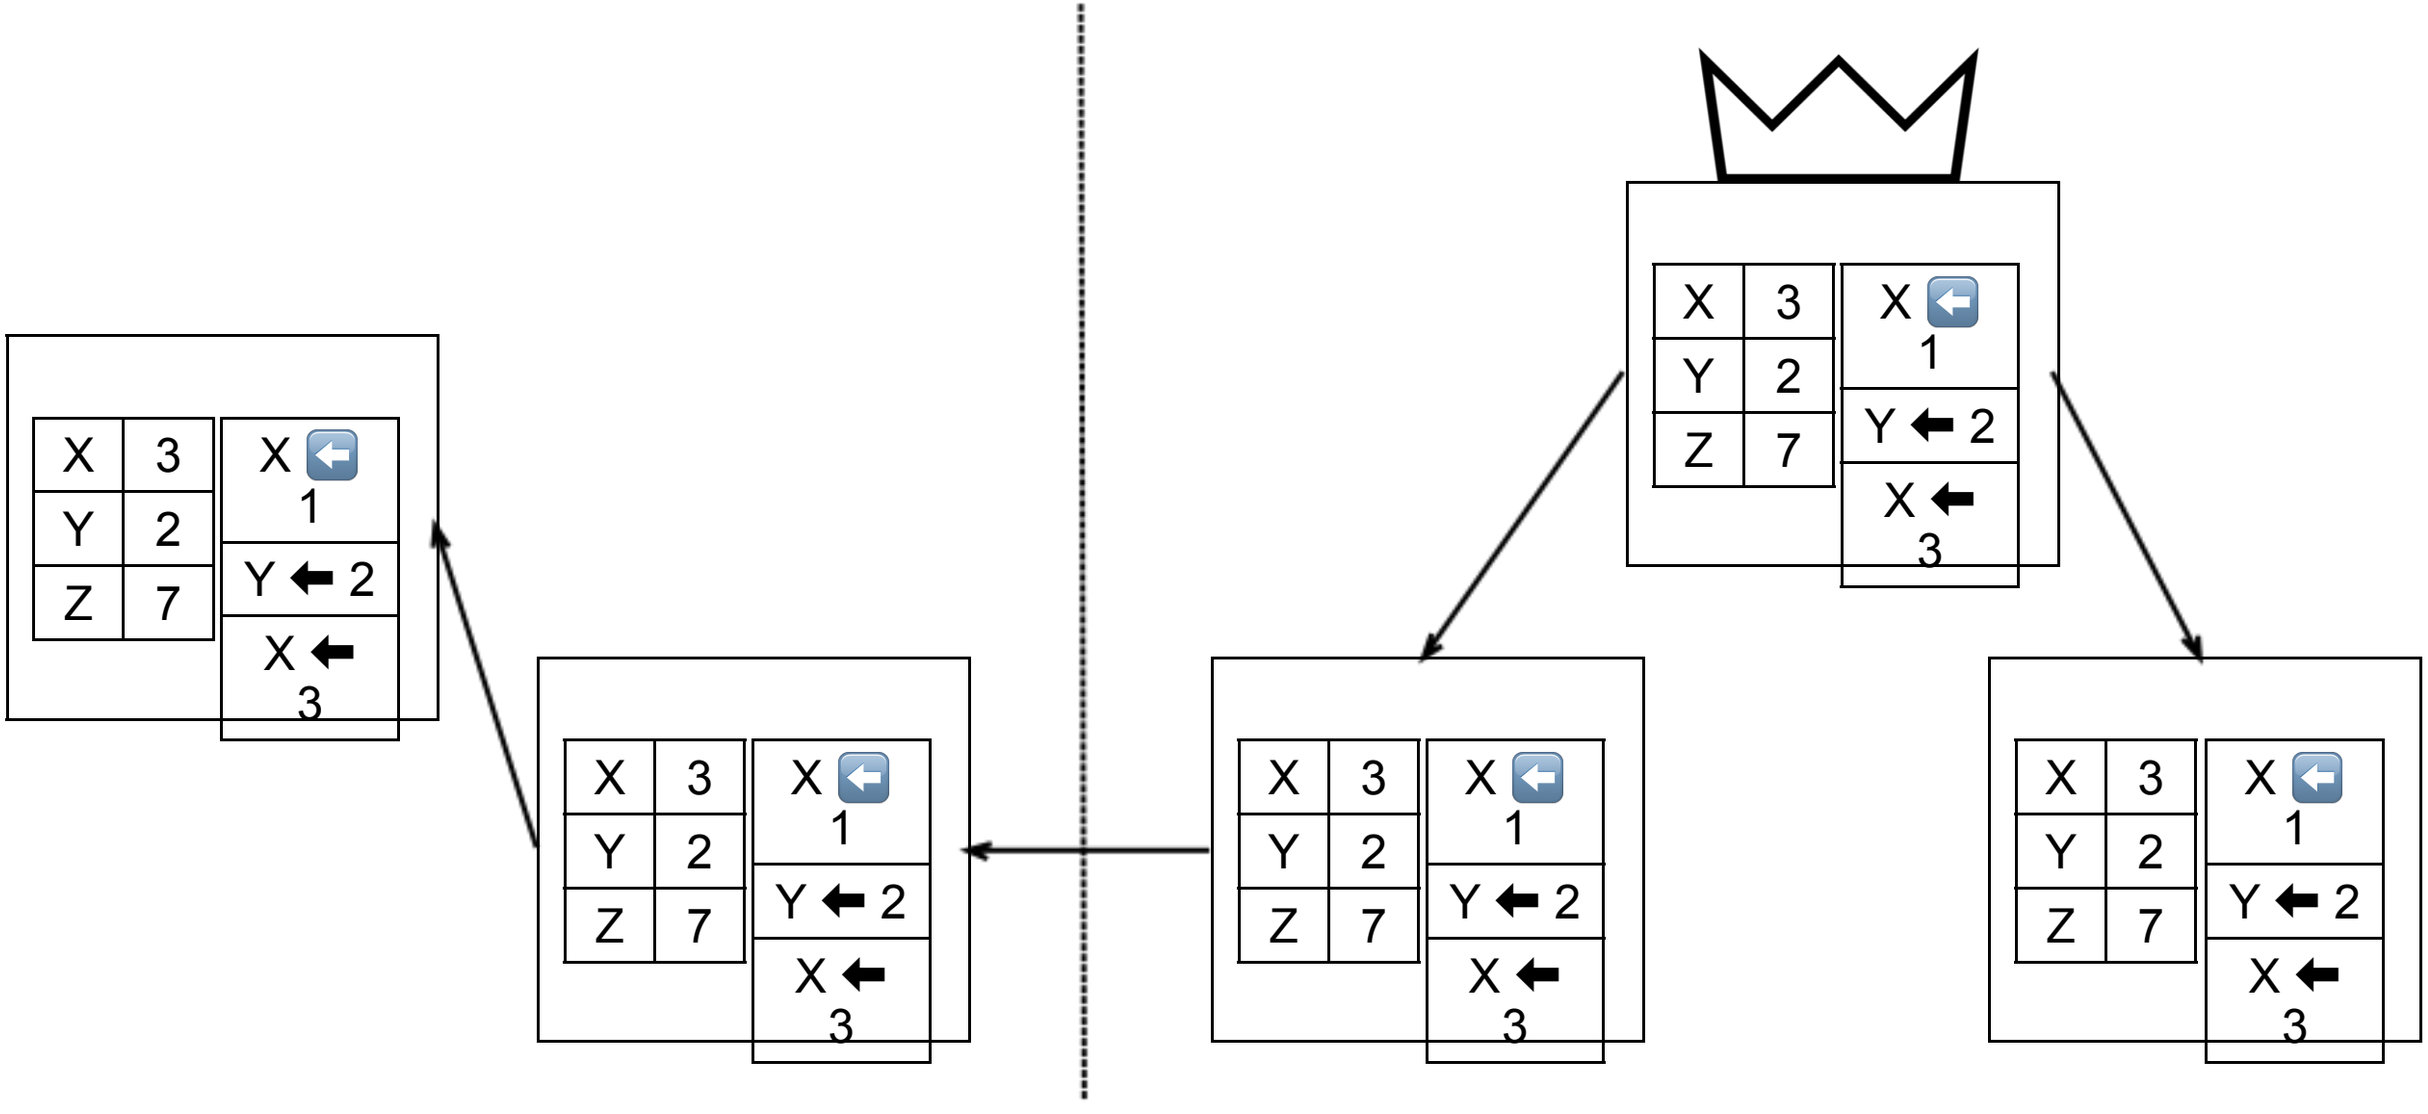

Determining When To Call For An Election

- Tradeoff between failover time and spurious failovers
- Node calls for an election when it hasn’t heard from the primary within the election timeout
- Starting in 3.2:
  - Election timeout is configurable
  - Election timeout is varied randomly for each node
  - Varying timeouts help reduce tied votes
  - Fewer tied votes = faster failover

Conclusion

- MongoDB 3.2 will have
  - Faster failovers
  - Faster error detection
  - More control to prevent spurious failovers
- This means your systems are
  - More stable
  - More resilient to failure
  - Easier to maintain

Questions?